

```
hive> create database company;
OK
Time taken: 0.57 seconds
hive> show databases;
OK
company
default
retail
Time taken: 0.115 seconds, Fetched: 3 row(s)
hive> use company;
OK
Time taken: 0.141 seconds
hive> create table company.employee(id int,name string,salary int,designation string,dept string);
OK
Time taken: 0.307 seconds

hive> insert into table employee values (1201,'Gopal',45000,'Technical Manager','TP');
Query ID = cloudera_20221114122828_887807d5-a85d-41b0-b327-58db7d8410f9
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1668453151595_0001, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1668453151595_0001/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1668453151595_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2022-11-14 12:29:27,803 Stage-1 map = 0%, reduce = 0%
2022-11-14 12:29:46,763 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.64 sec
MapReduce Total cumulative CPU time: 2 seconds 640 msec
Ended Job = job_1668453151595_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/company.db/employee/.hive-staging_hive_2022-11-14_12-28-32_357_4524765815586125032-1/-ext-10000
Loading data to table company.employee
Table company.employee stats: [numFiles=1, numRows=1, totalSize=38, rawDataSize=37]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Cumulative CPU: 2.64 sec HDFS Read: 4139 HDFS Write: 110 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 640 msec
OK
Time taken: 78.052 seconds
hive> select * from employee;
OK
1201 Gopal 45000 Technical Manager TP
Time taken: 0.266 seconds, Fetched: 1 row(s)
hive> insert into table employee values (1202,'Manisha',45000,'Proof Reader','PR'),
(1203,'Ram',40000,'Technical Writer','TP'),(1204,'Kiran',40000,'HR Admin','HR'),(1205,'Raj',30000,'OP
Admin','Admin');
Query ID = cloudera_20221114123232_77e6b111-8dcd-4155-b945-97551f6cd794
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1668453151595_0002, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1668453151595_0002/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1668453151595_0002
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2022-11-14 12:32:37,888 Stage-1 map = 0%, reduce = 0%
2022-11-14 12:33:00,516 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.42 sec
MapReduce Total cumulative CPU time: 2 seconds 420 msec
Ended Job = job_1668453151595_0002
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/company.db/employee/.hive-staging_hive_2022-11-14_12-32-20_502_3198761408279509719-1/-ext-10000
Loading data to table company.employee
Table company.employee stats: [numFiles=2, numRows=5, totalSize=167, rawDataSize=162]
```

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Cumulative CPU: 2.42 sec HDFS Read: 4318 HDFS Write: 202 SUCCESS

Total MapReduce CPU Time Spent: 2 seconds 420 msec

OK

Time taken: 43.132 seconds

```
select * from employee where salary>40000;
```

OK

1201	Gopal	45000	Technical Manager	TP
1202	Manisha	45000	Proof Reader	PR

```
hive> select * from employee where dept="TP";
```

OK

1201	Gopal	45000	Technical Manager	TP
1203	Ram	40000	Technical Writer	TP

Time taken: 0.349 seconds, Fetched: 2 row(s)

```
hive> select * from employee;
```

OK

1201	Gopal	45000	Technical Manager	TP
1202	Manisha	45000	Proof Reader	PR
1203	Ram	40000	Technical Writer	TP
1204	Kiran	40000	HR Admin	HR
1205	Raj	30000	OP Admin	Admin

Time taken: 0.391 seconds, Fetched: 5 row(s)